

# Databases

## ER Model

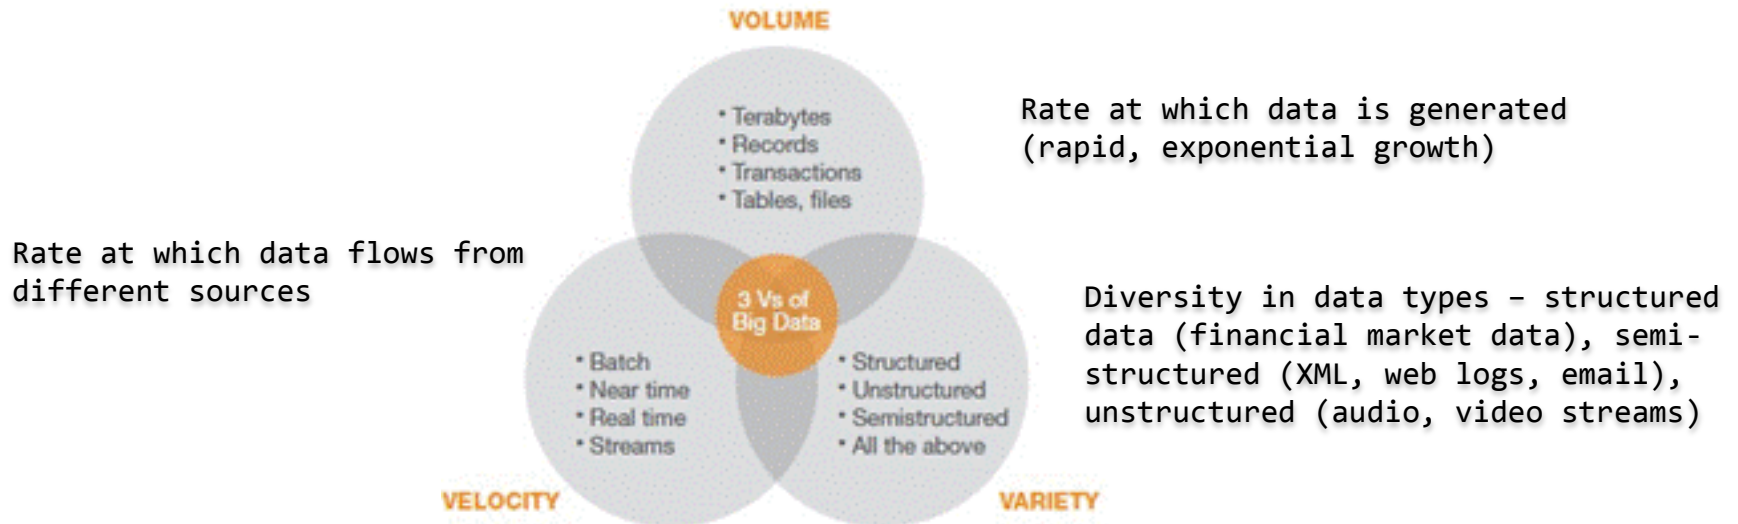
COMP 1531

Aarthi Natarajan

Week 09

# What is data?

- **Data** – facts that can be recorded and have implicit meaning (Elmasri & Navathe)
- Today data is being generated at an exponential rate
  - Financial market data, posts to social media sites, growing logs of web transactions, computation physics...BIG Data



# Why do we need a database ?

- Data by itself is not very useful.
- Give a context to data to transform data into **information** e.g., the numbers 45,55,67 do not mean much, but given a context such as these are the marks of students in COMP 1531, this is now information

**DATA -> INFORMATION -> DECISION**

- This data needs to be:
  - **Stored** (in a structured format)
  - **Manipulated** (efficiently, usefully)
  - **Shared** (by very many users, concurrently)
  - **Transmitted**
- Red text handled by **databases**; green by **networks**.

# Databases today...

- **Nearly every computer application uses a database**
  - Google, EBay, Amazon, iTunes Shop
  - Library catalogues, Train time tables, Airline bookings
  - Bank accounts, credit card, debit card
  - Medical records (Medicare), Tax Office
  - Facebook, Twitter, ...
- Every time you use a loyalty card, you're inputting information about your buying habits into the database of the company you are buying from
- **Challenges in building effective databases**
  - efficiency, security, scalability, maintainability, availability, integration, new media types (e.g. music), ...

# What is a database ?

- A **database** represents a logically coherent collection of related data
- A **database management system (DBMS)** is a software application that allows users to:
  - create and maintain a database (DDL)
  - define **queries** to retrieve data from the database
  - perform **transactions** that cause some data to be written or deleted from the database (DML)
  - provides concurrency, integrity, security to the database
- A database and DBMS are collectively referred to as a **database system**

# Data Models

A **data model** describes how the data is structured in the database

There are several types of data models

- **Relational model**
  - a data structure where data is stored as a set of records known as **tables**
  - each table consists of **rows** of information (also called a **tuple**)
  - each row contains fields known as **columns**
- **Document model**
  - data is stored in a hierarchical fashion e.g., XML
- **Object-oriented model**
  - a data structure where data is stored as a collection of objects
- **Object-Relational model**
  - a hybrid model that combines the relational and the object-oriented database models

StudentId	FirstName	LastName
213899	Joe	Bloggs
321456	Sam	Hunt
456789	John	Smith

# More database terminology

- A **database schema** adheres to a **data model** and provides a logical view of the database, defining how the data is organised and the relationships between them and is typically set up at the beginning
- A **database schema instance** is the state of the database at a particular instance of time


# Relational Database System

- A **Relational Database Management System** (RDBMS) is a DBMS that:
  - is based on a relational data model i.e., stores data as *tuples* or *records* in *tables*
  - Allows the user to create *relationships* between tables
- Examples of relational database systems
  - **Open Source**
    - PostgreSQL, MySQL, SQLite
  - **Commercial**
    - Oracle, DB2 (IBM), MS SQL Server, Sybase



# Database System Architecture

## Typical environment for a modern database system



SQL Queries and results travel along the client <-> server links

# **Data Modelling for Databases**

# Database design

Typical steps in a database design

1. requirements analysis (identify data and operations)
- 2. data modelling** (high-level, abstract)  
an important early stage of database application development (aka "database engineering")
3. database schema design (detailed, relational model/tables)
4. database implementation (create instance of schema)
5. build operations/interface (SQL, stored procedures, GUI)
6. performance tuning (physical re-design)
7. schema evolution (logical schema re-design)


# Data Modelling

Data modelling, in general consists of building:

- **Logical models**: abstract model e.g., ER Model, OO Model
- **Physical models**: record-based models e.g., relational model, classes which deal with the physical layout of data in storage

A **data-modelling strategy** for designing a database

- Design using abstract model (**conceptual-level modelling**)
  - i.e., **perform initial conceptual modelling with entity relationship (ER) models**
- Map to physical model (**implementation-level modelling**)
  - **Transform conceptual ER design into relational model**



# Data Modelling for Databases

## Aims of Data Modelling:

- describe what **data** is contained in the database  
(e.g. entities: students, courses, accounts, branches, patients, ...)
- describe **relationships** between data items  
(e.g. John is enrolled in COMP3311, Paul's account is held at Coogee)
- describe **constraints** on data  
(e.g. 7-digit IDs, students can enrol in no more than 30UC per semester)

Data modelling is a **design** process

- converts requirements into a data model

# Some Design Ideas

Consider the following during design:

- start simple ... evolve design as problem better understood
- identify objects (and their properties), then relationships
- most designs involve kinds (classes) of people
- keywords in requirements suggest data/relationships  
(rule-of-thumb: **nouns** → **data**, **verbs** → **relationships**)
- don't confuse operations with relationships  
(**operation**: he **buys** a book; **relationship**: the book **is owned** by him)
- consider all possible data, not just what's available

# Entity Relationship Diagrams

# Entity-Relationship Conceptual Data Modelling

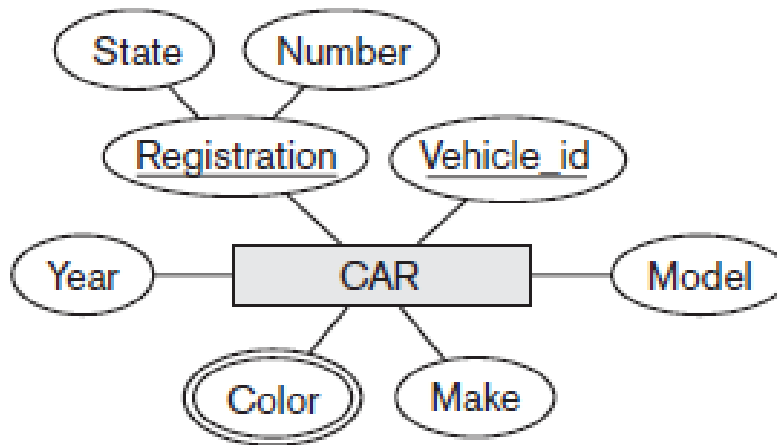
The world is viewed as a collection of **inter-related** entities.

ER modelling uses **three** major modelling constructs:

- **entity**:
  - a thing or object of interest in the real-world and is distinguishable from other objects
- **attribute**:
  - a data item or property of interest describing the entity  
*e.g., Joe (entity) described by name, address, age (attributes)*
- An **entity-set** (aka: entity-type) can be viewed as either:
  - a set of entities with the same set of attributes
  - an abstract description of a class of entities e.g., students, courses, accounts



e.g.,



An ER diagram showing an **entity-set** CAR with two **key attributes** (registration and vehicle\_id), three **single-valued attributes** (year, model, make) and a **multi-valued attribute** (color)

CAR  
Registration (Number, State), Vehicle\_id, Make, Model, Year, {Color}

CAR<sub>1</sub>  
((ABC 123, TEXAS), TK629, Ford Mustang, convertible, 2004 {red, black})

CAR<sub>2</sub>  
((ABC 123, NEW YORK), WP9872, Nissan Maxima, 4-door, 2005, {blue})

CAR<sub>3</sub>  
((VSY 720, TEXAS), TD729, Chrysler LeBaron, 4-door, 2002, {white, blue})

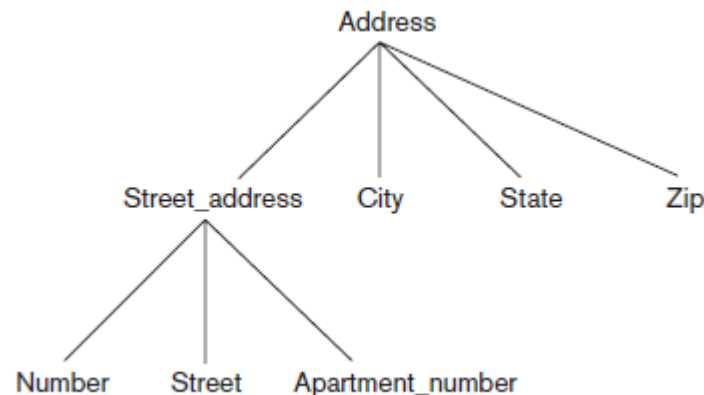
⋮

An **entity set** CAR with three entities

# Attributes of an entity-set

In contrast to relational model, attributes in an ER model can be:

- **Simple** (attribute cannot be broken into smaller sub-parts)
  - e.g., **age** attribute for entity type Employee
- **Composite** (have a hierarchy of attributes)
  - e.g., entity type EMPLOYEE has a composite attribute **Address**



- **Single-valued** (have only one value for each entity)
  - e.g., an **vin\_chassis** attribute for an entity type CAR
- **Multi-valued** ( have a set of values for each entity)
  - e.g., a **Colors** attribute for CAR = (blue,black)

# What if two entities have the same set of attribute values?

- They're regarded as the same entity.
- So, each entity must have a distinct set of attribute values.

## One approach:

Define a **key (super-key)** : It is any set of attributes

- whose set of values are distinct over entity set
- natural (e.g. **name + address + birthday**) or artificial (e.g. **SSN**)

- **Candidate key** = any super-key such that **no subset** is also a superkey)  
e.g. (**name + address**) is a super-key, but not (**name**) or (**address**)
- **Primary key** = a candidate key designated by DB designer that uniquely identifies an entity e.g., **SSN**

## Example (bank customer entities)

Customer = ( custNo, name, address, taxFileNo )

- Definite super-keys:
  - any set of attributes involving custNo or taxFileNo
- Possible super-keys:
  - ( name, address )
- Unlikely super-keys:
  - (name), (address)

# Relationship Sets

**Relationship**: relates two or more entities, e.g.,

- Joe Smith ( a STUDENT entity ) **ENROLLED IN** (relationship)  
COMP1531 ( a COURSE entity )
- Chris (an EMPLOYEE entity) **WORKS FOR** (relationship)  
ORACLE (a COMPANY entity)

**Relationship Set (aka relationship type)** : set of similar relationships, associating entities belonging to one entity-set to another

- **degree** = the number of entities involved in the relationship (in ER model,  $\geq 2$ ) e.g, the degree of **WORKS FOR** is 2
- **cardinality** = # associated entities on each side of relationship e.g., the cardinality of WORKS FOR is N:1

# ER model vs OO model

Analogy between ER and OO models:

- an **entity** is like an **object instance**
- an **entity set** is like a **class**

Differences between ER and OO models:

- ER modelling doesn't consider operations (methods)

# Entity Relationship Diagrams

- **ER diagrams** are a graphical tool for data modelling
- An ER diagram consists of:
  - a collection of *entity set* definitions
  - a collection of *relationship set* definitions
  - *attributes* associated with entity and relationship sets
  - connections between entity and relationship sets

**Warning:** 99% of the time ...

- we say "entity" when we mean "entity set"
- we say "relationship" when we mean "relationship set"
- If we want to refer to a specific entity, we generally say "entity instance"

# Entity Relationship Diagrams


Specific visual symbols indicate different ER design elements:




Entity



Weak entity




Relationship



Identifying Relationship



Attribute



Multi-valued attribute



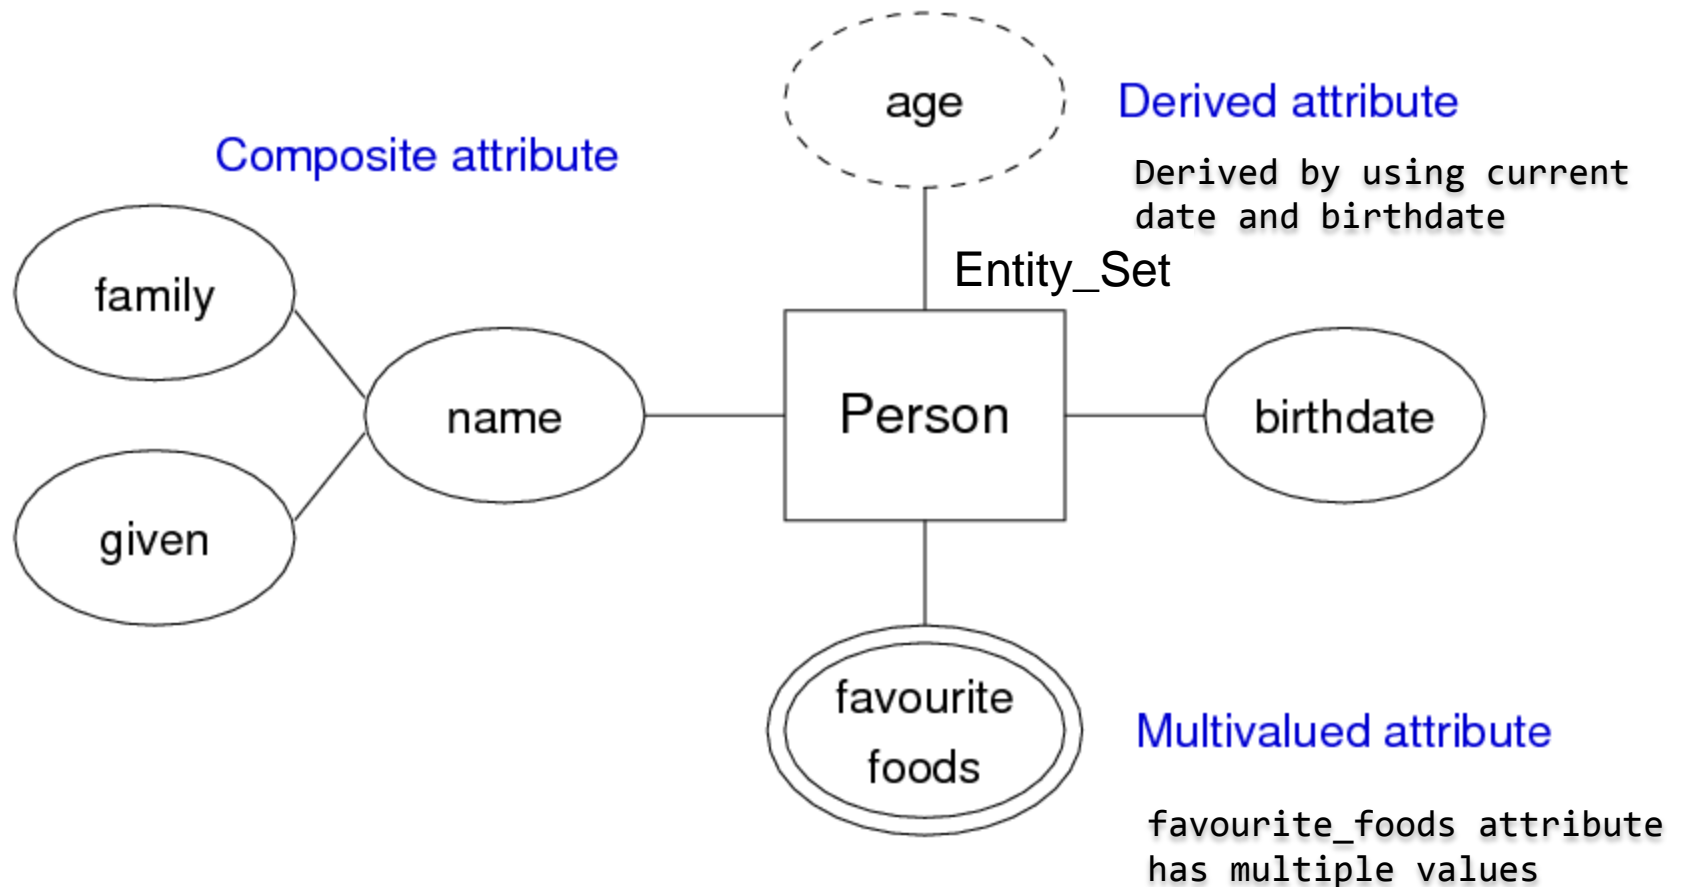
Inheritance



Derived attribute



# Example of attribute notations




# Cardinality in Relationship Sets

## Examples:


one-to-one



one-to-many




many-to-many



# An alternative explicit notation

## Examples:



# Relationship Sets in ER diagrams

**Level of participation constraint** = a type of relationship constraint defined as:

Participation in relationship set  $R$  by entity set  $A$  may be:

- **total** - every  $a \in A$  participates in  $\geq 1$  relationship in  $R$
- **partial** - only some  $a \in A$  participate in relationships in  $R$

Example:

- every bank loan is associated with at least one customer
- not every customer in a bank has a loan



# Exercise 1: Relationship Semantics

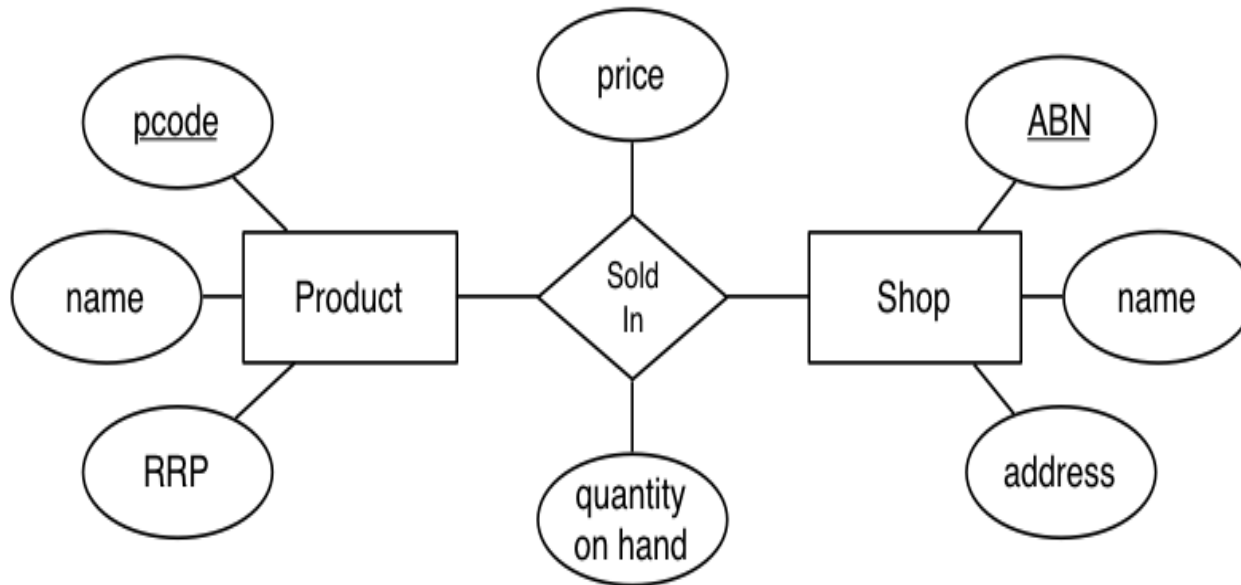
Describe precisely the scenarios implied by the following relationships:



# Relationship Type with attributes

In some cases, a relationship needs associated attributes

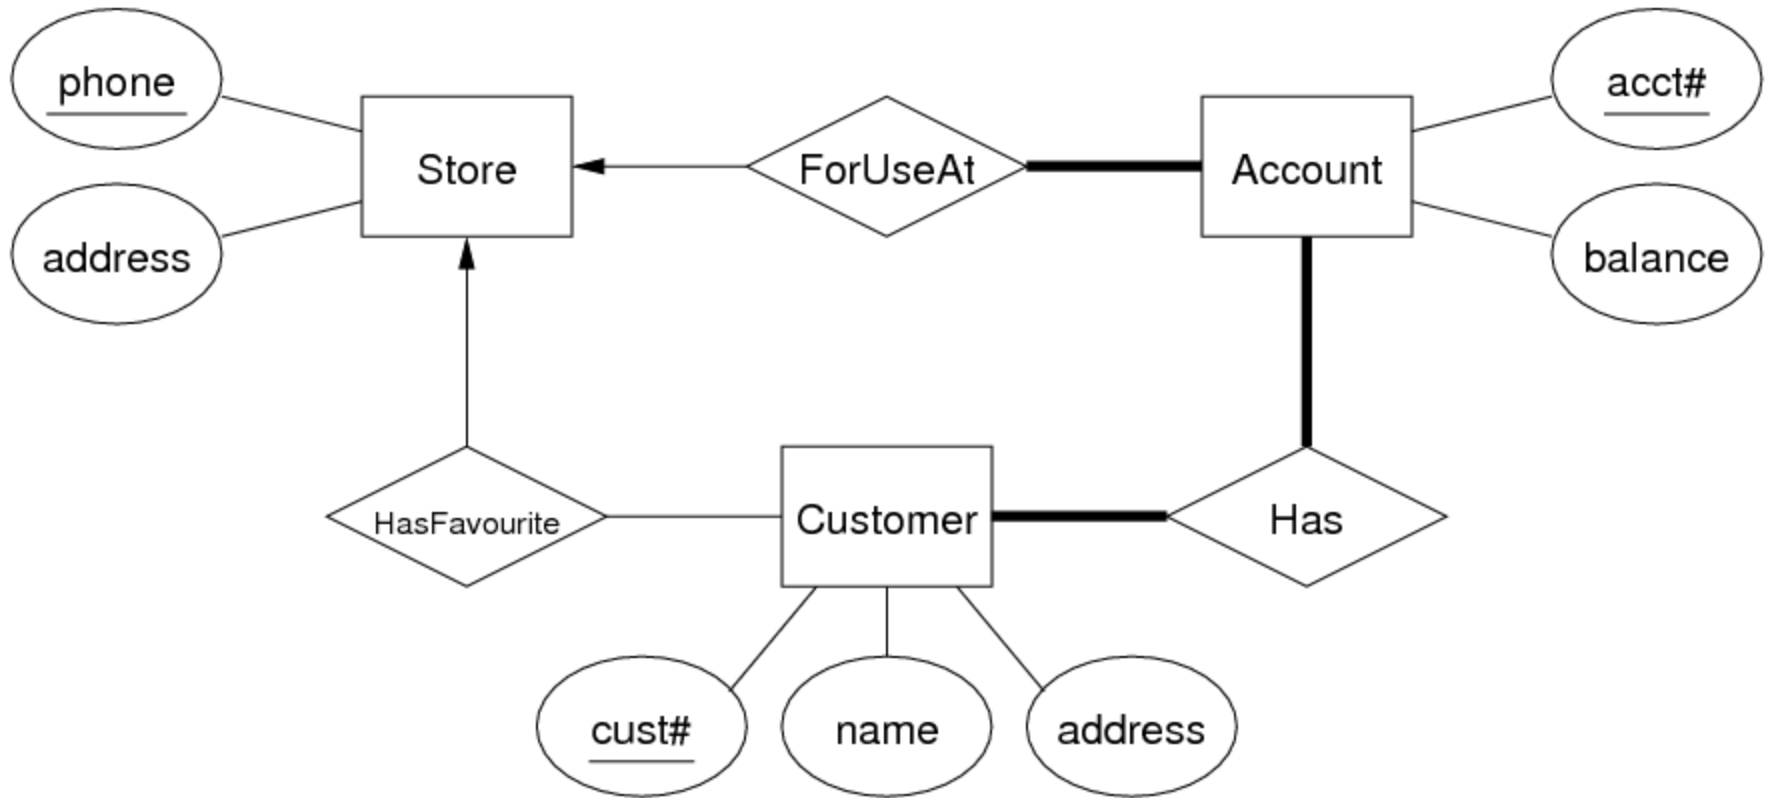
**Example:**



(price and quantity are related to products in a particular shop)

# Putting it all together...

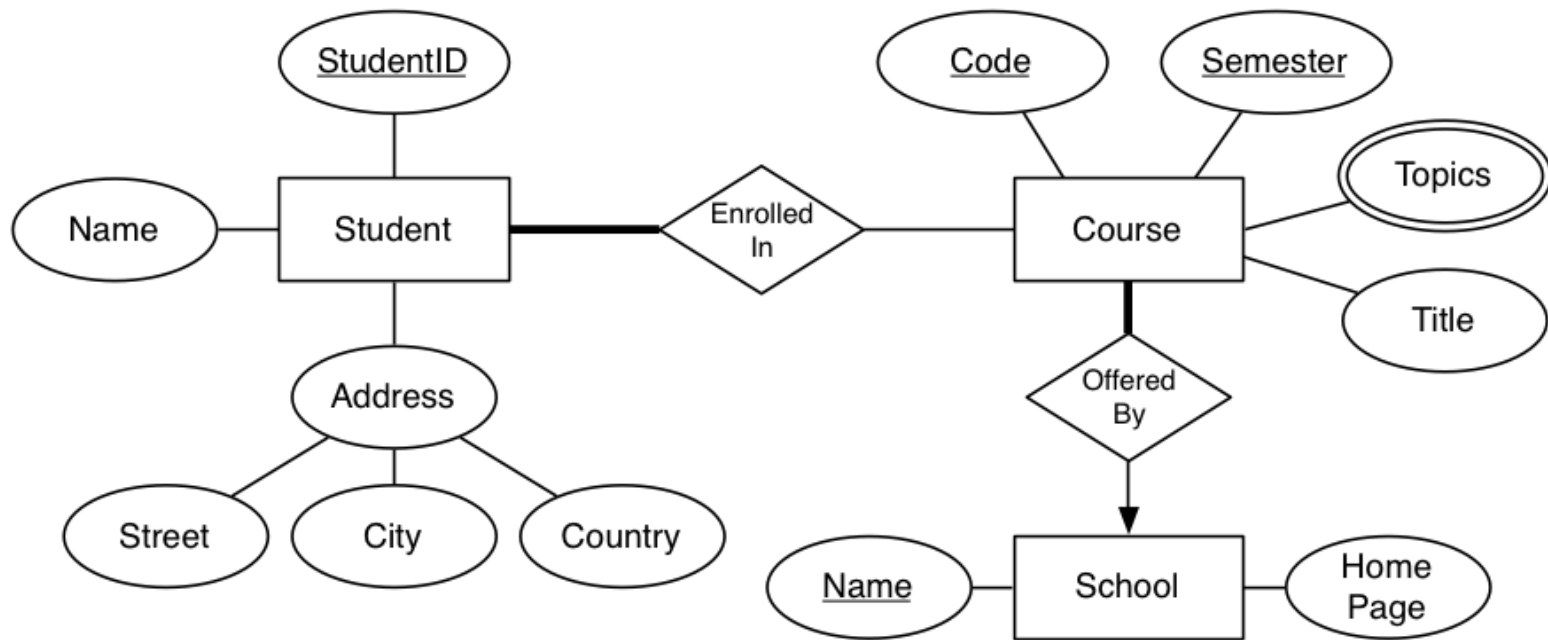
## Example1: - a complete ER Diagram



primary key attributes are underlined e.g. cust#

## Example 2:

Entities, relationships, attributes, keys, cardinality, participation, ...






# Weak Entity Set

A Weak entity set

- has no key of its own;
- exist only because of association with strong entities

Example:



# Subclasses and Inheritance

A **subclass** of an entity set  $A$  is a set of entities:

- with all attributes of  $A$ , plus (usually) it own attributes
- that is involved in all of  $A$ 's relationships, plus its own

Properties of subclasses:


- **overlapping** or **disjoint** (can an entity be in multiple subclasses?)
- **total** or **partial** (does every entity have to also be in a subclass?)

Special case: entity has one subclass ("B is-a A" specialisation)


# Subclasses and Inheritance

## Example:

*A person may be a doctor and/or may be a patient or may be neither*



*Every employee is either a permanent employee or works under a contract*



# Design considerations using the ER model

- should an "object" be represented by an attribute or entity?
- is a "concept" best expressed as an entity or relationship?
- should we use  $n$ -way rel<sup>n</sup>ship or several 2-way rel<sup>n</sup>ships?
- is an "object" a strong or weak entity? (usually strong)
- are there subclasses/superclasses within the entities?

Answers to above are worked out by *thinking* about the application domain.

## Exercise 1


Develop an ER design for the following scenario:

A database records information about employees and the departments they work for:

- For each employee, the name and emp\_id
- For each department, the name and allocated budget
- An employee may work for several departments for different periods of time
- A department may have several employees working for it

# Design considerations ...

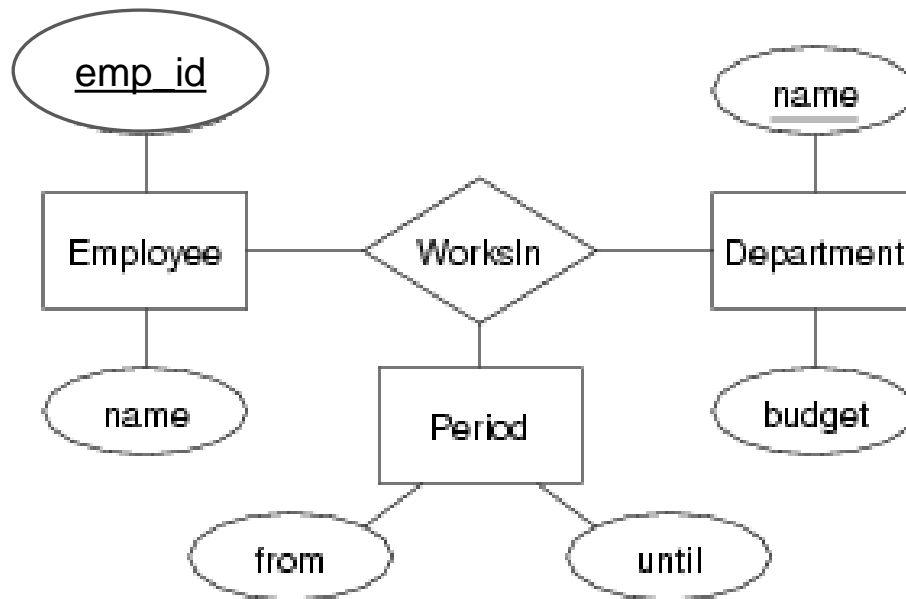
## Attribute vs Entity Example (v1)



**Assumption:** Employees can work for several departments, but cannot work for the same department over two different time periods.

# Design considerations ...

## Attribute vs Entity Example (v2)



**Assumption:** Employees can work for the same department over two different time periods.

# Design using the ER model

ER diagrams are typically too large to fit on a single screen.  
(or a single sheet of paper, if printing)


One commonly used strategy:

- define entity sets separately, showing attributes
- combine entities and relationships on a single diagram (but without showing entity attributes)
- if very large design, may use several linked diagrams as seen in the example in the next three set of slides




# e.g. an ER model for a Bank

## (1) Modelling people (employees)



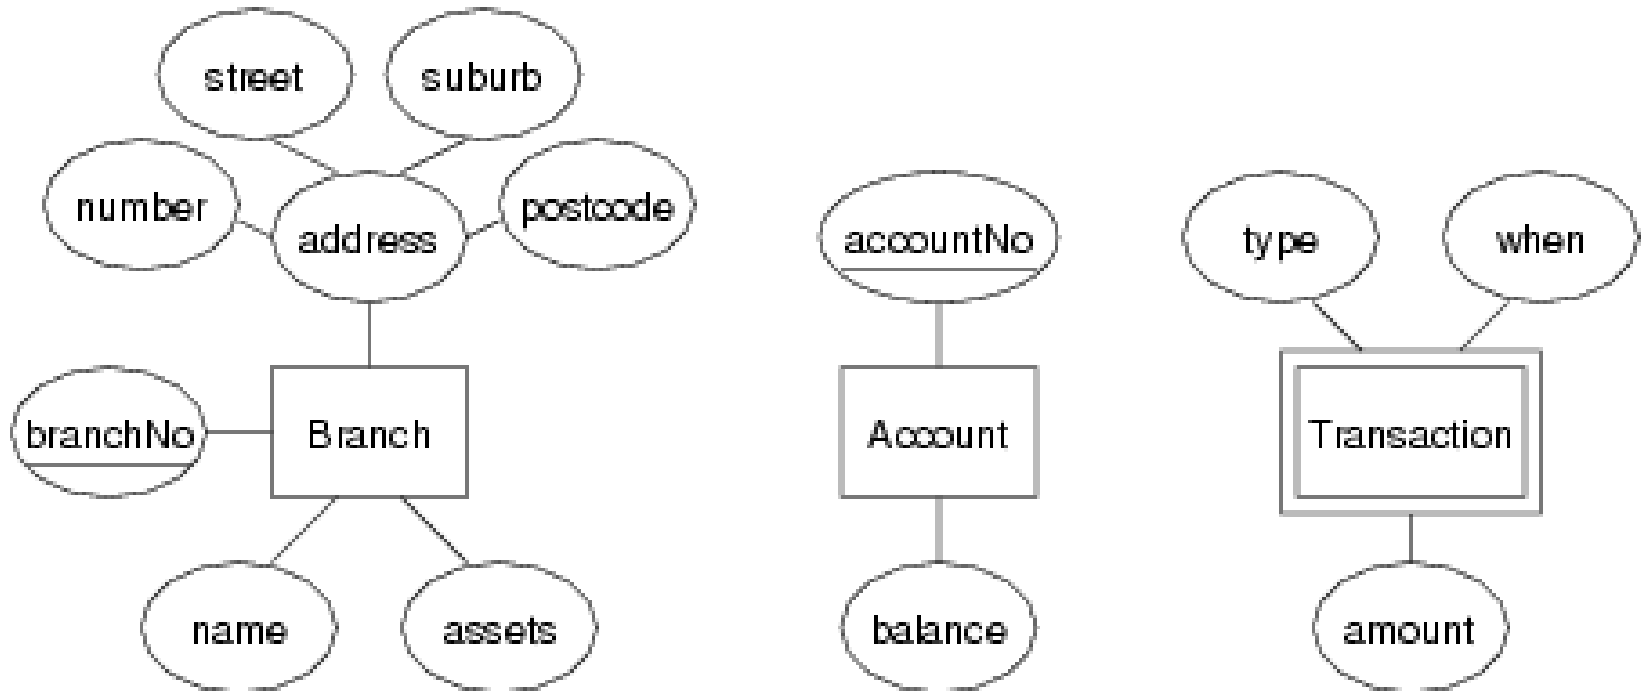
## (2) Modelling people (customer)

Modelling people (cont):



## e.g. an ER model for a Bank

### (3) Modelling branches, accounts, transactions



## e.g. an ER model for a Bank

### (4) Putting it all together with relationships

